



Workers Bindings

[Copy Page](#)

This guide will instruct you through setting up and deploying your first Workers AI project. You will use [Workers](#), a Workers AI binding, and a large language model (LLM) to deploy your first AI-powered application on the Cloudflare global network.

1. Sign up for a [Cloudflare account](#).
2. Install [Node.js](#).

Node.js version manager

1. Create a Worker project

You will create a new Worker project using the create-cloudflare CLI (C3). [C3](#) is a command-line tool designed to help you set up and deploy new applications to Cloudflare.

Create a new project named hello-ai by running:

npm yarn pnpm bun

```
npm create cloudflare@latest -- hello-ai
```

Running `npm create cloudflare@latest` will prompt you to install the [create-cloudflare package](#), and lead you through setup. C3 will also install [Wrangler](#), the Cloudflare Developer Platform CLI.

For setup, select the following options:

- For *What would you like to start with?*, choose `Hello World Starter`.
- For *Which template would you like to use?*, choose `Worker only`.
- For *Which language do you want to use?*, choose `TypeScript`.
- For *Do you want to use git for version control?*, choose `Yes`.
- For *Do you want to deploy your application?*, choose `No` (we will be making some changes before deploying).

This will create a new `hello-ai` directory. Your new `hello-ai` directory will include:

- A "Hello World" [Worker](#) at `src/index.ts`.
- A [wrangler.jsonc](#) configuration file.

Go to your application directory:

```
cd hello-ai
```

2. Connect your Worker to Workers AI

You must create an AI binding for your Worker to connect to Workers AI. [Bindings](#) allow your Workers to interact with resources, like Workers AI, on the Cloudflare Developer Platform.

To bind Workers AI to your Worker, add the following to the end of your Wrangler file:

`{}` wrangler.jsonc  `wrangler.toml`

```
{
  "ai": {
    "binding": "AI"
  }
}
```

Your binding is [available in your Worker code](#) on `env.AI`.

You can also bind Workers AI to a Pages Function. For more information, refer to [Functions](#)

[Bindings](#).

3. Run an inference task in your Worker

You are now ready to run an inference task in your Worker. In this case, you will use an LLM, [llama-3.1-8b-instruct](#), to answer a question.

Update the `index.ts` file in your `hello-ai` application directory with the following code:

JS JavaScript

TS TypeScript

index.js

```
export default {
  async fetch(request, env) {
    const response = await env.AI.run("@cf/meta/llama-3.1-8b-instruct", {
      prompt: "What is the origin of the phrase Hello, World",
    });

    return new Response(JSON.stringify(response));
  },
};
```

Up to this point, you have created an AI binding for your Worker and configured your Worker to be able to execute the Llama 3.1 model. You can now test your project locally before you deploy globally.

4. Develop locally with Wrangler

While in your project directory, test Workers AI locally by running [wrangler dev](#):

```
npx wrangler dev
```

Workers AI local development usage charges

Using Workers AI always accesses your Cloudflare account in order to run AI models and will incur usage charges even in local development.

You will be prompted to log in after you run the `wrangler dev`. When you run `npx wrangler dev`, Wrangler will give you a URL (most likely `localhost:8787`) to review your Worker. After you go to the URL Wrangler provides, a message will render that resembles the following example:

```
{
  "response": "Ah, a most excellent question, my dear human friend! *adjusts
               glasses*\n\nThe origin of the phrase \"Hello, World\" is a fascinating tale
               that spans several decades and multiple disciplines. It all began in the
               early days of computer programming, when a young man named Brian Kernighan
               was tasked with writing a simple program to demonstrate the basics of a new
               programming language called C.\nKernighan, a renowned computer scientist and
               author, was working at Bell Labs in the late 1970s when he created the
               program. He wanted to showcase the language's simplicity and versatility, so
               he wrote a basic \"Hello, World!\" program that printed the familiar greeting
               to the console.\n\nThe program was included in Kernighan and Ritchie's
               influential book \"The C Programming Language,\" published in 1978. The book
               became a standard reference for C programmers, and the \"Hello, World!\"
               program became a sort of \"Hello, World!\" for the programming community.
               \nOver time, the phrase \"Hello, World!\" became a shorthand for any simple
               program that demonstrated the basics"
}
```

5. Deploy your AI Worker

Before deploying your AI Worker globally, log in with your Cloudflare account by running:

```
npx wrangler login
```

You will be directed to a web page asking you to log in to the Cloudflare dashboard. After you have logged in, you will be asked if Wrangler can make changes to your Cloudflare account. Scroll down and select **Allow** to continue.

Finally, deploy your Worker to make your project accessible on the Internet. To deploy your Worker, run:

```
npx wrangler deploy
```

```
https://hello-ai.<YOUR_SUBDOMAIN>.workers.dev
```

Your Worker will be deployed to your custom [workers.dev](#) subdomain. You can now visit the URL to run your AI Worker.

By finishing this tutorial, you have created a Worker, connected it to Workers AI through an AI binding, and ran an inference task from the Llama 3 model.

Related resources

- [Cloudflare Developers community on Discord](#) ↗ - Submit feature requests, report bugs, and share your feedback directly with the Cloudflare team by joining the Cloudflare Discord server.
- [Models](#) - Browse the Workers AI models catalog.
- [AI SDK](#) - Learn how to integrate with an AI model.

Was this helpful?



Previous

← Getting started

Next

REST API →

Last updated: Apr 3, 2025

Resources

[API](#)

[New to Cloudflare?](#)

[Products](#)

[Sponsorships](#)

[Open Source](#)

Company

[cloudflare.com](#)

[Our team](#)

[Careers](#)

Community

[!\[\]\(104fbf564e2e5a8fbd84f31656d114c7_img.jpg\) X](#)

[!\[\]\(aab88c0d099e5d18d6533a97b13ec28d_img.jpg\) Discord](#)

[!\[\]\(b538fe54c1f3a7343e37e85cc2d00497_img.jpg\) YouTube](#)

[!\[\]\(5abce1a84a655b073239ab33e1199487_img.jpg\) GitHub](#)

Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

Tools

[Cloudflare Radar](#)

[Speed Test](#)

[Is BGP Safe Yet?](#)

[RPKI Toolkit](#)

[Certificate Transparency](#)

2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#) • [Report Security Issues](#) • [Trademark](#)

• [Your Privacy Choices](#)