



qwen2.5-coder-32b-instruct

Text Generation • Qwen

@cf/qwen/qwen2.5-coder-32b-instruct

Qwen2.5-Coder is the latest series of Code-Specific Qwen large language models (formerly known as CodeQwen). As of now, Qwen2.5-Coder has covered six mainstream model sizes, 0.5, 1.5, 3, 7, 14, 32 billion parameters, to meet the needs of different developers. Qwen2.5-Coder brings the following improvements upon CodeQwen1.5:

Model Info

Context Window ↗	32,768 tokens
LoRA	Yes
Unit Pricing	\$0.66 per M input tokens, \$1.00 per M output tokens

Playground

Try out this model with Workers AI LLM Playground. It does not require any setup or authentication and an instant way to preview and test a model directly in the browser.

Launch the LLM Playground

Usage

Worker - Streaming

Worker

Python

```
import os
import requests

ACCOUNT_ID = "your-account-id"
AUTH_TOKEN = os.environ.get("CLOUDFLARE_AUTH_TOKEN")

prompt = "Tell me all about PEP-8"
response = requests.post(
    f"https://api.cloudflare.com/client/v4/accounts/{ACCOUNT_ID}/ai/run/@cf/qwen/qwen2.5-coder-32b-instruct",
    headers={"Authorization": f"Bearer {AUTH_TOKEN}"},
    json={
        "messages": [
            {"role": "system", "content": "You are a friendly assistant"},
            {"role": "user", "content": prompt}
        ]
    }
)
result = response.json()
print(result)
```

curl

OpenAI compatible endpoints

Workers AI also supports OpenAI compatible API endpoints for `/v1/chat/completions` and `/v1/embeddings`. For more details, refer to [Configurations](#).

Parameters

* indicates a required field

Input

0 **object**

prompt * **string** min 1

The input text prompt for the model to generate a response.

lora **string**

Name of the LoRA (Low-Rank Adaptation) model to fine-tune the base model.

response_format **object**

type **string**

json_schema

raw **boolean**

If true, a chat template is not applied and you must adhere to the specific model's expected formatting.

stream **boolean**

If true, the response will be streamed back incrementally using SSE, Server Sent Events.

max_tokens **integer** default 256

The maximum number of tokens to generate in the response.

temperature **number** default 0.6 min 0 max 5

Controls the randomness of the output; higher values produce more random results.

top_p **number** min 0 max 2

Adjusts the creativity of the AI's responses by controlling how many possible words it considers. Lower values make outputs more predictable; higher values allow for more varied and creative responses.

top_k **integer** min 1 max 50

Limits the AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises.

seed **integer** min 1 max 9999999999

Random seed for reproducibility of the generation.

repetition_penalty **number** min 0 max 2

Penalty for repeated tokens; higher values discourage repetition.

frequency_penalty **number** min 0 max 2

Decreases the likelihood of the model repeating the same lines verbatim.

presence_penalty **number** min 0 max 2

Increases the likelihood of the model introducing new topics.

1 **object**

messages * **array**

An array of message objects representing the conversation history.

items **object**

role * **string**

The role of the message sender (e.g., 'user', 'assistant', 'system', 'tool').

content * **string**

The content of the message as a string.

functions **array**

items **object**

name * **string**

code * **string**

tools **array**

A list of tools available for the assistant to use.

items **one of**

 0 **object**

name * **string**

The name of the tool. More descriptive the better.

description * **string**

A brief description of what the tool does.

parameters * **object**

Schema defining the parameters accepted by the tool.

type * **string**

The type of the parameters object (usually 'object').

required **array**

List of required parameter names.

items **string**

properties * **object**

Definitions of each parameter.

additionalProperties **object**

type * **string**

The data type of the parameter.

description * **string**

A description of the expected parameter.

1 **object**

type * **string**

Specifies the type of tool (e.g., 'function').

function * **object**

Details of the function tool.

name * **string**

The name of the function.

description * **string**

A brief description of what the function does.

parameters * **object**

Schema defining the parameters accepted by the function.

type * **string**

The type of the parameters object (usually 'object').

required **array**

List of required parameter names.

items **string**

properties * **object**

Definitions of each parameter.

additionalProperties **object**

type * **string**

The data type of the parameter.

description * **string**

A description of the expected parameter.

response_format **object**

type **string**

json_schema

raw **boolean**

If true, a chat template is not applied and you must adhere to the specific model's expected formatting.

stream **boolean**

If true, the response will be streamed back incrementally using SSE, Server Sent Events.

max_tokens **integer** default 256

The maximum number of tokens to generate in the response.

temperature **number** default 0.6 min 0 max 5

Controls the randomness of the output; higher values produce more random results.

top_p **number** min 0 max 2

Adjusts the creativity of the AI's responses by controlling how many possible words it considers. Lower values make outputs more predictable; higher values allow for more varied and creative responses.

top_k **integer** min 1 max 50

Limits the AI to choose from the top 'k' most probable words. Lower values make responses more focused; higher values introduce more variety and potential surprises.

seed **integer** min 1 max 9999999999

Random seed for reproducibility of the generation.

repetition_penalty **number** min 0 max 2

Penalty for repeated tokens; higher values discourage repetition.

frequency_penalty **number** min 0 max 2

Decreases the likelihood of the model repeating the same lines verbatim.

presence_penalty **number** min 0 max 2

Increases the likelihood of the model introducing new topics.

Output

0 **object**

response * **string**

The generated text response from the model

usage **object**

Usage statistics for the inference request

prompt_tokens **number** 0

Total number of tokens in input

completion_tokens **number** 0

Total number of tokens in output

total_tokens **number** 0

Total number of input and output tokens

tool_calls **array**

An array of tool calls requests made during the response generation

items **object**

arguments **object**

The arguments passed to be passed to the tool call request

name **string**

The name of the tool to be called

1 **string**

API Schemas

The following schemas are based on JSON Schema

Input

Output

```
{
  "type": "object",
  "oneOf": [
    {
      "title": "Qwen2_5_Coder_32B_Instruct_Prompt",
      "properties": {
        "prompt": {
          "type": "string",
          "minLength": 1,
          "description": "The input text prompt for the model to
generate a response."
        },
        "lora": {
          "type": "string",
          "description": "Name of the LoRA (Low-Rank Adaptation)
model to fine-tune the base model."
        },
        "response_format": {
          "title": "JSON Mode",
          "type": "object",
          "properties": {
            "type": {
              "type": "string",
              "enum": [
                "json_object",
                "json_schema"
              ]
            },
            "json_schema": {}
          }
        },
        "raw": {
          "type": "boolean",
          "default": false,
          "description": "If true, a chat template is not applied
and you must adhere to the specific model's expected formatting."
        },
        "stream": {
          "type": "boolean".
```

```
    "type": "boolean",
    "default": false,
    "description": "If true, the response will be streamed
back incrementally using SSE, Server Sent Events."
  },
  "max_tokens": {
    "type": "integer",
    "default": 256,
    "description": "The maximum number of tokens to generate
in the response."
  },
  "temperature": {
    "type": "number",
    "default": 0.6,
    "minimum": 0,
    "maximum": 5,
    "description": "Controls the randomness of the output;
higher values produce more random results."
  },
  "top_p": {
    "type": "number",
    "minimum": 0,
    "maximum": 2,
    "description": "Adjusts the creativity of the AI's
responses by controlling how many possible words it considers. Lower values
make outputs more predictable; higher values allow for more varied and
creative responses."
  },
  "top_k": {
    "type": "integer",
    "minimum": 1,
    "maximum": 50,
    "description": "Limits the AI to choose from the top 'k'
most probable words. Lower values make responses more focused; higher values
introduce more variety and potential surprises."
  },
  "seed": {
    "type": "integer",
    "minimum": 1,
    "maximum": 9999999999,
```

```
        "description": "Random seed for reproducibility of the
generation."
    },
    "repetition_penalty": {
        "type": "number",
        "minimum": 0,
        "maximum": 2,
        "description": "Penalty for repeated tokens; higher
values discourage repetition."
    },
    "frequency_penalty": {
        "type": "number",
        "minimum": 0,
        "maximum": 2,
        "description": "Decreases the likelihood of the model
repeating the same lines verbatim."
    },
    "presence_penalty": {
        "type": "number",
        "minimum": 0,
        "maximum": 2,
        "description": "Increases the likelihood of the model
introducing new topics."
    }
},
"required": [
    "prompt"
]
},
{
    "title": "Qwen2_5_Coder_32B_Instruct_Messages",
    "properties": {
        "messages": {
            "type": "array",
            "description": "An array of message objects representing
the conversation history.",
            "items": {
                "type": "object",
                "properties": {
```

```
        "role": {
          "type": "string",
          "description": "The role of the message
sender (e.g., 'user', 'assistant', 'system', 'tool')."
        },
        "content": {
          "type": "string",
          "description": "The content of the message as
a string."
        }
      },
      "required": [
        "role",
        "content"
      ]
    },
    "functions": {
      "type": "array",
      "items": {
        "type": "object",
        "properties": {
          "name": {
            "type": "string"
          },
          "code": {
            "type": "string"
          }
        },
        "required": [
          "name",
          "code"
        ]
      }
    },
    "tools": {
      "type": "array",
      "description": "A list of tools available for the
assistant to use.",
```

```
    "items": {
      "type": "object",
      "oneOf": [
        {
          "properties": {
            "name": {
              "type": "string",
              "description": "The name of the tool.
More descriptive the better."
            },
            "description": {
              "type": "string",
              "description": "A brief description
of what the tool does."
            },
            "parameters": {
              "type": "object",
              "description": "Schema defining the
parameters accepted by the tool.",
              "properties": {
                "type": {
                  "type": "string",
                  "description": "The type of
the parameters object (usually 'object')."
                },
                "required": {
                  "type": "array",
                  "description": "List of
required parameter names.",
                  "items": {
                    "type": "string"
                  }
                },
                "properties": {
                  "type": "object",
                  "description": "Definitions
of each parameter.",
                  "additionalProperties": {
                    "type": "object",
                    "description": "A brief description
of what the tool does."
                  }
                }
              }
            }
          }
        }
      ]
    }
  }
}
```

"The data type of the parameter."

description of the expected parameter."

```

        "properties": {
          "type": {
            "type": "string",
            "description":
              "The data type of the parameter."
          },
          "description": {
            "type": "string",
            "description": "A
              description of the expected parameter."
          }
        },
        "required": [
          "type",
          "description"
        ]
      }
    },
    "required": [
      "type",
      "properties"
    ]
  }
},
"required": [
  "name",
  "description",
  "parameters"
]
},
{
  "properties": {
    "type": {
      "type": "string",
      "description": "Specifies the type of
        tool (e.g., 'function')."
    },
    "function": {
      "type": "object"
    }
  }
}

```

```
function tool.",
the function."

description of what the function does."

defining the parameters accepted by the function.",
type of the parameters object (usually 'object')."

of required parameter names.",

"Definitions of each parameter.",

"additionalProperties": {
    "type": "object",
    "properties": {
```



```
        "type": {
          "type":
"string",

"description": "The data type of the parameter."

},

"description": {

"type":
"string",

"description": "A description of the expected parameter."
}
},
"required": [
  "type",
  "description"
]
}
},
"required": [
  "type",
  "properties"
]
}
},
"required": [
  "name",
  "description",
  "parameters"
]
}
},
"required": [
  "type",
  "function"
]
}
```

```
    ]
  }
},
"response_format": {
  "title": "JSON Mode",
  "type": "object",
  "properties": {
    "type": {
      "type": "string",
      "enum": [
        "json_object",
        "json_schema"
      ]
    },
    "json_schema": {}
  }
},
"raw": {
  "type": "boolean",
  "default": false,
  "description": "If true, a chat template is not applied
and you must adhere to the specific model's expected formatting."
},
"stream": {
  "type": "boolean",
  "default": false,
  "description": "If true, the response will be streamed
back incrementally using SSE, Server Sent Events."
},
"max_tokens": {
  "type": "integer",
  "default": 256,
  "description": "The maximum number of tokens to generate
in the response."
},
"temperature": {
  "type": "number",
  "default": 0.6,
  "minimum": 0,
```

```
        "maximum": 5,
        "description": "Controls the randomness of the output;
higher values produce more random results."
    },
    "top_p": {
        "type": "number",
        "minimum": 0,
        "maximum": 2,
        "description": "Adjusts the creativity of the AI's
responses by controlling how many possible words it considers. Lower values
make outputs more predictable; higher values allow for more varied and
creative responses."
    },
    "top_k": {
        "type": "integer",
        "minimum": 1,
        "maximum": 50,
        "description": "Limits the AI to choose from the top 'k'
most probable words. Lower values make responses more focused; higher values
introduce more variety and potential surprises."
    },
    "seed": {
        "type": "integer",
        "minimum": 1,
        "maximum": 9999999999,
        "description": "Random seed for reproducibility of the
generation."
    },
    "repetition_penalty": {
        "type": "number",
        "minimum": 0,
        "maximum": 2,
        "description": "Penalty for repeated tokens; higher
values discourage repetition."
    },
    "frequency_penalty": {
        "type": "number",
        "minimum": 0,
        "maximum": 2,
        ... ..
```

```
        "description": "Decreases the likelihood of the model
repeating the same lines verbatim."
    },
    "presence_penalty": {
        "type": "number",
        "minimum": 0,
        "maximum": 2,
        "description": "Increases the likelihood of the model
introducing new topics."
    }
},
"required": [
    "messages"
]
}
]
```

Was this helpful?



[Previous](#)

[← Dashboard](#)

[Next](#)

[Workers Bindings →](#)

Resources

[API](#)

[New to Cloudflare?](#)

[Products](#)

[Sponsorships](#)

[Open Source](#)

Support

[Help Center](#)

[System Status](#)

[Compliance](#)

[GDPR](#)

Company

[cloudflare.com](#)

[Our team](#)

[Careers](#)

Tools

[Cloudflare Radar](#)

[Speed Test](#)

[Is BGP Safe Yet?](#)

[RPKI Toolkit](#)

[Certificate Transparency](#)

Community

[!\[\]\(0fb13ad0bfa3d86868cdd3883e5665b3_img.jpg\) X](#)

[!\[\]\(799877f5c2f906134441300079881630_img.jpg\) Discord](#)

[!\[\]\(41aea2746216b27a6939d696d8e035da_img.jpg\) YouTube](#)

[!\[\]\(7bc43b319a082987e20f7bf78f4bab80_img.jpg\) GitHub](#)

2025 Cloudflare, Inc. • [Privacy Policy](#) • [Terms of Use](#) • [Report Security Issues](#) • [Trademark](#)

• [Your Privacy Choices](#)